

## Car Price Prediction using Random Forest Regressor

Create a New Code Snippet □ ↑ ↓ ± ⌂

### Project Overview

The aim of this Data Science Project to focuses on building a machine learning model to predict car price for a company in Poland. Identifying car price by their model, generation, year of production, mileage, localization, and type and volume of engine.

```
: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt
```

#### • Variables:

Price: Price variable shows the cost of car in Poland currency PLN (approx. 1USD=1PLN)  
Unnamed: 0: It represent the serial number.  
mark: Mark of the car.  
model: Model of the car.  
generation\_name: Formatted Generation Name of the car.  
year: Car Year of formatted.  
mileage: Car Mileage in Kilometres.  
vol\_engine: Auto Engine Size.  
fuel: Represent Engine Type according to their fuel type.  
city: Locality in Poland.  
province: Region of Poland.

```
: df = pd.read_csv("Car_Prices.csv")  
df.head()
```

Activate Windows

Go to Settings to activate Wind

	Unnamed: 0	mark	model	generation_name	year	mileage	vol_engine	fuel	city	province	price
0	0	opel	combo	gen-d-2011	2015	139568	1248	Diesel	Janki	Mazowieckie	35900
1	1	opel	combo	gen-d-2011	2018	31991	1499	Diesel	Katowice	Śląskie	78501
2	2	opel	combo	gen-d-2011	2015	278437	1598	Diesel	Brzeg	Opolskie	27000
3	3	opel	combo	gen-d-2011	2016	47600	1248	Diesel	Korfantów	Opolskie	30800
4	4	opel	combo	gen-d-2011	2014	103000	1400	CNG	Tarnowskie Góry	Śląskie	35900

### drop irravalent feature

```
: df.drop("Unnamed: 0",axis=1, inplace=True)
```

```
: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 117927 entries, 0 to 117926  
Data columns (total 10 columns):  
 #   Column      Non-Null Count  Dtype    
 ---    
 0   mark        117927 non-null  object   
 1   model       117927 non-null  object   
 2   generation_name  87842 non-null  object   
 3   year         117927 non-null  int64   
 4   year         117927 non-null  int64   
 5   mileage      117927 non-null  int64   
 6   vol_engine   117927 non-null  int64   
 7   fuel          117927 non-null  object   
 8   city          117927 non-null  object   
 9   province     117927 non-null  object  
dtypes: int64(4), object(6)  
memory usage: 9.0+ MB
```

Activate Windows

Go to Settings to activate Wind

### drop all the null and duplicated values

```
df.dropna(inplace=True)
```

```
df.duplicated().sum()
```

5046

```
#drop all the duplicated values  
df.drop_duplicates(inplace=True)
```

```
df.describe().astype(int)
```

	year	mileage	vol_engine	price
count	82796	82796	82796	82796

Activate Windows

Go to Settings to activate Wind

	2012	152751	1842	62598
mean	5	89996	613	75624
std	1978	0	0	900
min	2008	86338	1498	19900
25%	2012	158700	1798	37900
50%	2017	212316	1995	73850
75%	2022	2500000	6600	2399900
max				

```
#for numeric columns
columns = df.select_dtypes(include='number')
```

```
columns
```

	year	mileage	vol_engine	price
0	2015	139568	1248	35900
1	2018	31991	1499	78501
2	2015	278437	1598	27000
3	2016	47600	1248	30800
4	2014	103000	1400	35900
...	...	...	...	...
117250	2017	51000	1969	229900
117251	2016	83500	1969	135000
117252	2017	174000	1969	154500
117253	2016	189020	1969	130000
117254	2016	95000	1969	126000

82796 rows × 4 columns

```
import seaborn as sns

fig, axes= plt.subplots(2,2, figsize=(10,5))

sns.histplot(data=df, x='year', kde=True, color="blue", edgecolor="black",ax=axes[0,0])
axes[0, 0].set_title('Distribution of Year')



sns.histplot(data=df, x='mileage', kde=True, color="green", edgecolor="black",ax=axes[0,1])
axes[0, 1].set_title('Distribution of Mileage')

sns.histplot(data=df, x='vol_engine', kde=True, color="green", edgecolor="black",ax=axes[1,0])
axes[1, 0].set_title('Distribution of Engine Volume')


sns.histplot(data=df, x='price', kde=True, color="red", edgecolor="black",ax=axes[1,1])
axes[1, 1].set_title('Distribution of Price')

plt.tight_layout()

plt.show()
```



Activate Windows  
Go to Settings to activate Wind



```
def remove_outliers(df, column):
    Q1 = df[column].quantile(0.25)
    Q3 = df[column].quantile(0.75)

    IQR = Q3 - Q1

    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    df_filtered = df[(df[column] >= lower_bound) & (df[column] <= upper_bound)]

    num_outliers = len(df) - len(df_filtered)
    print(f"\n--- Outlier Removal Summary for '{column}' ---")
    print(f"Q1: {Q1:.2f}, Q3: {Q3:.2f}, IQR: {IQR:.2f}")
    print(f"Lower Bound: {lower_bound:.2f}, Upper Bound: {upper_bound:.2f}")
    print(f"Original rows: {len(df)}, Filtered rows: {len(df_filtered)}")
    print(f"Number of outliers removed: {num_outliers}")

    print(f"Original rows: {len(df)}, Filtered rows: {len(df_filtered)}")
    print(f"Number of outliers removed: {num_outliers}")

    return df_filtered
```

```
for col in columns:
    df = remove_outliers(df, col)
```

```
--- Outlier Removal Summary for 'year' ---
Q1: 2008.00, Q3: 2017.00, IQR: 9.00
Lower Bound: 1994.50, Upper Bound: 2030.50
Original rows: 82796, Filtered rows: 82645
Number of outliers removed: 151
```

```
--- Outlier Removal Summary for 'mileage' ---
Q1: 86180.00, Q3: 212021.00, IQR: 125841.00
Lower Bound: -102581.50, Upper Bound: 400782.50
Original rows: 82645, Filtered rows: 82228
Number of outliers removed: 417
```

```
--- Outlier Removal Summary for 'vol_engine' ---
Q1: 1498.00, Q3: 1995.00, IQR: 497.00
Lower Bound: 752.50, Upper Bound: 2740.50
Original rows: 82228, Filtered rows: 75046
Number of outliers removed: 7182
```

```
--- Outlier Removal Summary for 'price' ---
Q1: 19500.00, Q3: 68899.00, IQR: 49399.00
Lower Bound: -54598.50, Upper Bound: 142997.50
Original rows: 75046, Filtered rows: 69488
```

```
--- Outlier Removal Summary for 'price' ---
Q1: 19500.00, Q3: 68899.00, IQR: 49399.00
Lower Bound: -54598.50, Upper Bound: 142997.50
Original rows: 75046, Filtered rows: 69488
Number of outliers removed: 5558
```

Activate Windows  
Go to Settings to activate Wind

### After removed the outlier

```
[15]: fig, axes = plt.subplots(2,2, figsize=(10,5))

sns.histplot(data=df, x='year', kde=True, color="blue", edgecolor="black", ax=axes[0,0])
axes[0, 0].set_title('Distribution of Year')


sns.histplot(data=df, x='mileage', kde=True, color="green", edgecolor="black", ax=axes[0,1])
axes[0, 1].set_title('Distribution of Mileage')

sns.histplot(data=df, x='vol_engine', kde=True, color="green", edgecolor="black", ax=axes[1,0])
axes[1, 0].set_title('Distribution of Engine Volume')

sns.histplot(data=df, x='price', kde=True, color="red", edgecolor="black", ax=axes[1,1])
axes[1, 1].set_title('Distribution of Price')

plt.tight_layout()
plt.show()
```

Activate Windows  
Go to Settings to activate Wind



```
df.describe()
```

	year	mileage	vol_engine	price
<b>count</b>	69488.000000	69488.000000	69488.000000	69488.000000
<b>mean</b>	2011.855961	159203.062716	1686.756131	42062.148788
<b>std</b>	5.141381	79948.075522	337.784631	30519.505297
<b>min</b>	1995.000000	1.000000	875.000000	900.000000
<b>25%</b>	2008.000000	102392.500000	1399.000000	18800.000000
<b>50%</b>	2012.000000	165000.000000	1598.000000	32900.000000
<b>75%</b>	2016.000000	212900.000000	1984.000000	59600.000000
<b>max</b>	2022.000000	400500.000000	2720.000000	142990.000000




```
# to getting text carrying columns
text_columns = df.select_dtypes(include='object')
```


	mark	model	generation_name	fuel	city	province
0	opel	combo	gen-d-2011	Diesel	Janki	Mazowieckie
1	opel	combo	gen-d-2011	Diesel	Katowice	Śląskie
3	opel	combo	gen-d-2011	Diesel	Korfantów	Opolskie
4	opel	combo	gen-d-2011	CNG	Tarnowskie Góry	Śląskie
...	...	...	...	...	...	...
117230	volvo	xc-90	gen-ii-2014-xc-90	Diesel	Wrocław	Dolnośląskie
117237	volvo	xc-90	gen-ii-2014-xc-90	Gasoline	Warszawa	Mazowieckie
117251	volvo	xc-90	gen-ii-2014-xc-90	Gasoline	Pruszcz Gdański	Pomorskie
117253	volvo	xc-90	gen-ii-2014-xc-90	Gasoline	Sionna	Mazowieckie
117254	volvo	xc-90	gen-ii-2014-xc-90	Gasoline	Warszawa	Mazowieckie

69488 rows × 6 columns

To know the Quantity of objects feature



```
for col in text_columns:
    plt.figure(figsize = (8,3))
    df[col].value_counts(ascending=False).head(10).plot(kind="bar")
    plt.show()
```







### To check the relation of features with the price


```
n [129]:  
for col in text_columns:  
    plt.figure(figsize = (8,4))  
    df.groupby(col)["price"].mean().sort_values(ascending=False).head(40).plot(kind="bar")  
    plt.show()
```







Activate Windows  
Go to Settings to activate Wind



Activate Window:  
Go to Settings to activa

```

from sklearn.preprocessing import LabelEncoder, StandardScaler
scaler = StandardScaler()
le = LabelEncoder()

```

to assign the numeric values to text columns

```

for col in df.select_dtypes(include="object"):
    df[col] = le.fit_transform(df[col])

```

```
df.corr()
```

	mark	model	generation_name	year	mileage	vol_engine	fuel	city	province	price
mark	1.000000	0.196819	0.246185	0.121745	-0.117552	-0.227961	0.052916	0.005299	0.001786	0.022738
model	0.196819	1.000000	-0.083408	0.071234	-0.018935	0.064960	-0.076213	-0.004996	-0.001162	0.123390
generation_name	0.246185	-0.083408	1.000000	0.240800	-0.239778	-0.253888	0.073241	0.007266	0.019613	0.100035
year	0.121745	0.071234	0.240800	1.000000	-0.671815	-0.250313	-0.077639	0.018455	0.004956	0.796587
mileage	-0.117552	-0.018935	-0.239778	-0.671815	1.000000	0.427870	-0.198991	-0.020707	-0.025703	-0.590134
vol_engine	-0.227961	0.064960	-0.253888	-0.250313	0.427870	1.000000	-0.300485	0.003427	-0.040087	0.077804
fuel	0.052916	-0.076213	0.073241	-0.077639	-0.198991	-0.300485	1.000000	0.008100	0.007001	-0.083026
city	0.005299	-0.004996	0.007266	0.018455	-0.020707	0.003427	0.008100	1.000000	-0.070049	0.026105
province	0.001786	-0.001162	0.019613	0.004956	-0.025703	-0.040087	0.007001	-0.070049	1.000000	-0.012935
price	0.022738	0.123390	0.100035	0.796587	-0.590134	0.077804	-0.083026	0.026105	-0.012935	1.000000

```
df.columns
```

```

Index(['mark', 'model', 'generation_name', 'year', 'mileage', 'vol_engine',
       'fuel', 'city', 'province', 'price'],
      dtype='object')

```

select features for training the models

```

x = df[['mark','model','generation_name','year','mileage','vol_engine','fuel']]
y = df['price']

```

```
#convert the values in the standard form
x = scaler.fit_transform(x)
```

```
from sklearn.model_selection import train_test_split,GridSearchCV
```

```
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.25,random_state=42)
```

```
from sklearn.metrics import mean_absolute_error,mean_squared_error,r2_score,mean_absolute_percentage_error
```

```
# write the function for predicting the result easily
def result(predict):
    print(f"mean_absolute_error is { mean_absolute_error(y_test,predict)}")
    print(f"mean_squared_error is { np.sqrt(mean_squared_error(y_test,predict))}")
    print(f"r2_score is { r2_score(y_test,predict)}")
    print(f"mean_absolute_percentage_error(y_test,predict)}")
```

```
from sklearn.tree import DecisionTreeRegressor
dtr = DecisionTreeRegressor()
```

Activate V  
Go to Setting

Activate V  
Go to Setting

```
: dtr.fit(x_train, y_train)
:   ▾ DecisionTreeRegressor ⓘ ⓘ
DecisionTreeRegressor()

: dtr_pred = dtr.predict(x_test)

: result(dtr_pred)
mean_absolute_error is 6349.8935625805
mean_squared_error is 10303.527483857199
r2_score is 0.883536882477968
0.18629390851940947

: from sklearn.neighbors import KNeighborsRegressor
knr = KNeighborsRegressor()

: knr.fit(x_train,y_train)
:   ▾ KNeighborsRegressor ⓘ ⓘ
KNeighborsRegressor()

: knr_pred = knr.predict(x_test)

: result(knr_pred)
mean_absolute_error is 5352.348457287589
mean_squared_error is 8463.709915392914
r2_score is 0.9214152848250615
0.15901738520827738

: from sklearn.ensemble import RandomForestRegressor
rfr = RandomForestRegressor()

grid_rfr = {
    'random_state': [None, 42],
    'n_estimators': [100,110,120],
}

gsv=GridSearchCV(rfr,grid_rfr, cv=5)

gsv.fit(x_train,y_train)
:   ▾ GridSearchCV ⓘ ⓘ
:   ▾ best_estimator_: RandomForestRegressor
    ▾ RandomForestRegressor ⓘ
:   ▾ best_params_
{'n_estimators': 120, 'random_state': 42}

grid_rfr_pred = gsv.predict(x_test)

result(grid_rfr_pred)
mean_absolute_error is 5040.714346060589
mean_squared_error is 8040.00448751887
r2_score is 0.9299864665432411
0.1485926172699571

gsv.best_params_
{'n_estimators': 120, 'random_state': 42}

import joblib

joblib.dump(scaler,"scaler")

['scaler']

joblib.dump(gsv,"rfr car price predictor model.pkl")
['rfr car price predictor model.pkl']
```